



Binary Search Trees


Chapter 11






Binary Search

- Binary search can perform nearest neighbor queries on an ordered map that is implemented with an array, sorted by key
 - at each step, the number of candidate items is halved
 - terminates after $O(\log n)$ steps
- Example: **find**(7)



Binary Search Trees


- A binary search tree is a binary tree storing keys (or key-value entries) at its internal nodes and satisfying the following property:
 - Let u , v , and w be three nodes such that u is in the left subtree of v and w is in the right subtree of v . We have $key(u) \leq key(v) \leq key(w)$
- External nodes do not store items
- An inorder traversal of a binary search tree visits the keys in increasing order



Search



- ❑ To search for a key k , we trace a downward path starting at the root
- ❑ The next node visited depends on the comparison of k with the key of the current node
- ❑ If we reach a leaf, the key is not found
- ❑ Example: **get(4)**:
 - Call `TreeSearch(4, root)`
- ❑ The algorithms for nearest neighbor queries are similar

```
Algorithm TreeSearch( $k, v$ )  
  if  $T.isExternal(v)$   
    return  $v$   
  if  $k < key(v)$   
    return TreeSearch( $k, left(v)$ )  
  else if  $k = key(v)$   
    return  $v$   
  else {  $k > key(v)$  }  
    return TreeSearch( $k, right(v)$ )
```





Insertion

- ❑ To perform operation **put**(k, o), we search for key k (using TreeSearch)
- ❑ Assume k is not already in the tree, and let w be the leaf reached by the search
- ❑ We insert k at node w and expand w into an internal node
- ❑ Example: insert 5





Deletion

- ❑ To perform operation **remove(k)**, we search for key k
- ❑ Assume key k is in the tree, and let v be the node storing k
- ❑ If node v has a leaf child w , we remove v and w from the tree with operation **removeExternal(w)**, which removes w and its parent
- ❑ Example: remove 4





Deletion (cont.)

- We consider the case where the key k to be removed is stored at a node v whose children are both internal
 - we find the internal node w that follows v in an inorder traversal
 - we copy $key(w)$ into node v
 - we remove node w and its left child z (which must be a leaf) by means of operation **removeExternal**(z)
- Example: remove 3




Deletion (cont.)

- We consider the case where the key k to be removed is stored at a node v whose children are both internal
 - we find the internal node w that precedes v in an inorder traversal
 - we copy $key(w)$ into node v
 - we remove node w and its right child z (which must be a leaf) by means of operation **removeExternal**(z)
- Example: remove 8




Performance

- Consider an ordered map with n items implemented by means of a binary search tree of height h
 - the space used is $O(n)$
 - methods **get**, **put** and **remove** take $O(h)$ time
- The height h is $O(n)$ in the worst case and $O(\log n)$ in the best case
- On average, a binary tree generated from random insertions and removals of keys has expected height $O(\log n)$




AVL Trees



AVL Tree Definition

- ❑ Inventors - Adel'son, Vel'skii and Landis
- ❑ AVL trees are balanced
- ❑ An AVL Tree is a **binary search tree** such that for every internal node v of T , the heights of the children of v can differ by at most 1
- ❑ **Height-balance property**




An example of an AVL tree where the heights are shown next to the nodes

Height of an AVL Tree

Fact: The **height** of an AVL tree storing n keys is $O(\log n)$.

Proof (by induction): Let us bound $m(h)$: the minimum number of internal nodes of an AVL tree of height h .

- We easily see that $m(1) = 1$ and $m(2) = 2$
- For $n > 1$, an AVL tree of height h contains the root node, one AVL subtree of height $h-1$ and another of height $h-2$ (or $h-1$, but will we pick $h-1$?).
- That is, $m(h) = 1 + m(h-1) + m(h-2)$
- Knowing $m(h-1) \geq m(h-2)$, we get $m(h) > 2m(h-2)$. So
 $m(h) > 2m(h-2)$, $m(h) > 4m(h-4)$, $m(h) > 8m(h-6)$, ... (by induction),
 $m(h) > 2^i m(h-2i)$
- Solving the base case we get: $m(h) > 2^{h-1/2}$
- Taking logarithms: $h < 2\log m(h) + 2$
- Thus, the height of an AVL tree is $O(\log n)$



A Sharper Bound

- Induction on minimum number of nodes in an AVL tree of height h :
 $m(h) \geq c^{h-1}$
- Base case: $h = 1, m(h) \geq 1$
- Suppose the claim is true for all $h < k$
- We have to show that $m(k) \geq c^{k-1}$
- Recall the recurrence relation:
 - $m(k) = m(k-1) + m(k-2) + 1$
 $\geq c^{k-2} + c^{k-3}$
- We will be able to show that $m(k) \geq c^{k-1}$ if we can show that $c^{k-2} + c^{k-3} \geq c^{k-1}$
- So c should be such that $c^2 - c - 1 \leq 0$
- The quadratic equation $c^2 - c - 1 \leq 0$ has roots $\frac{1-\sqrt{5}}{2}$ and $\frac{1+\sqrt{5}}{2}$
- Hence, we can take c as $\frac{1+\sqrt{5}}{2}$, which is approximately 1.63

Hence, an AVL tree with n nodes has height at most $\log_{1.63} n$

Structure of AVL Tree

- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.

Structure of AVL Tree

- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.
 - All nodes at levels $1, 2, \dots, k-2$ have 2 children.

Structure of AVL Tree

- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.
 - All nodes at levels $1, 2, \dots, k-2$ have 2 children. (proof by contradiction)
 - ◆ all levels 1 to $k-1$ are full (all nodes are present)
 - ◆ Hence the tree has at least 2^{k-1} nodes

Structure of AVL Tree

- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.
 - All nodes at levels $1, 2, \dots, k-2$ have 2 children. (proof by contradiction)
 - ◆ all levels 1 to $k-1$ are full (all nodes are present)
 - ◆ Hence the tree has at least 2^{k-1} nodes

Structure of AVL Tree

- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.
 - ◆ hence tree has at most 2^{2k-1} nodes.
 - All nodes at levels $1, 2, \dots, k-2$ have 2 children.
 - ◆ all levels 1 to $k-1$ are full
 - ◆ hence tree has at least 2^{k-1} nodes

Structure of AVL Tree

- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.
 - ◆ hence tree has at most 2^{2k-1} nodes.
 - All nodes at levels $1, 2, \dots, k-2$ have 2 children.
 - ◆ all levels 1 to $k-1$ are full
 - ◆ hence tree has at least 2^{k-1} nodes
 - Thus $2^{k-1} \leq n \leq 2^{2k-1}$

Structure of AVL Tree


- Consider an AVL tree on n nodes and the leaf that is closest to the root is at level k . Then
 - The height of the tree is at most $2k-1$.
 - ◆ hence tree has at most 2^{2k-1} nodes.
 - All nodes at levels $1, 2, \dots, k-2$ have 2 children.
 - ◆ all levels 1 to $k-1$ are full
 - ◆ hence tree has at least 2^{k-1} nodes
 - Thus $2^{k-1} \leq n \leq 2^{2k-1}$
- Substituting $h = 2k-1$, we get $2^{(h-1)/2} \leq n \leq 2^h$

Summary - Structure of AVL Tree


- In an AVL tree of height h , the leaf closest to the root is at level $(h+1)/2$
- On the first $(h-1)/2$ levels, the AVL tree is a complete binary tree
 - thins out after $(h-1)/2$ level
- $2^{(h-1)/2} \leq \text{number of nodes} \leq 2^h$

Insertion

- ❑ Insertion is as in a binary search tree
- ❑ Always done by expanding an external node.
- ❑ Example: insert 54




before insertion




after insertion

Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.




before insertion




after insertion

Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.




before insertion




after insertion

Insertion

- ❑ Inserting a node into an AVL tree changes the height of some of the nodes in T.
- ❑ The only nodes whose heights can increase are the ancestors of inserted node.




before insertion



after insertion


Insertion

- ❑ Inserting a node into an AVL tree changes the height of some of the nodes in T.
- ❑ The only nodes whose heights can increase are the ancestors of inserted node.
- ❑ If the insertion causes T to become unbalanced, then some ancestor of the inserted node would have a height imbalance.





Insertion

- ❑ Inserting a node into an AVL tree changes the height of some of the nodes in T .
- ❑ The only nodes whose heights can increase are the ancestors of inserted node.
- ❑ If the insertion causes T to become unbalanced, then some ancestor of the inserted node would have a height imbalance.
- ❑ We travel up the tree from the inserted node (v), until we find the first node (x) such that its grandparent (z) is unbalanced.
- ❑ Let y be the parent of node x .




Insertion

- To rebalance the subtree rooted at z, we must perform a **rotation**.




Rotations

- Rotation is a way of locally reorganizing a BST.
- Let x, y be two nodes such that $y = \text{parent}(x)$.
- $\text{Keys}(T_1) < \text{Key}(x) < \text{Keys}(T_2) < \text{Key}(y) < \text{Keys}(T_3)$




Rotations

- Rotation is a way of locally reorganizing a BST.
- Let x, y be two nodes such that $y = \text{parent}(x)$.
- $\text{Keys}(T_1) < \text{Key}(x) < \text{Keys}(T_2) < \text{Key}(y) < \text{Keys}(T_3)$



Restructuring – Single Rotation


- Rotations to make y the top-most node



- Insertion happens in subtree T_2
- $H(T_2)$ increases from h to $h+1$
- x remains balanced
 - $H(T_3) = h$ or $h+1$ or $h+2$
 - ◆ $h+2$ – then x is originally unbalanced
 - ◆ $h+1$ – then height of x does not increase – z is balanced
 - ◆ therefore $H(T_3) = h$
- $H(x)$ increases from $h+1$ to $h+2$


Restructuring – Single Rotation

- Rotations to make y the top-most node
- Insertion happens in subtree T_2
- y remains balanced
 - $H(T_1) = h+1$ or $h+2$ or $h+3$
 - ◆ $h+3$ – y is originally unbalanced
 - ◆ $h+2$ – height of y does not increase – z is balanced
 - ◆ therefore $H(T_1) = h+1$
- $H(y) = h+2$ to $h+3$.



Restructuring – Single Rotation


- Rotations to make y the top-most node



- Insertion happens in subtree T_2
- y remains balanced
 - $H(T_1) = h+1$ or $h+2$ or $h+3$
 - ♦ $h+3$ – y is originally unbalanced
 - ♦ $h+2$ – height of y does not increase – z is balanced
 - ♦ therefore $H(T_1) = h+1$
- $H(y) = h+2$ to $h+3$.
- z is unbalanced
 - $H(T_0) = h+1$ or $h+2$ or $h+3$
 - ♦ Since originally z was balanced
 - ♦ $H(T_0) = h+1$
- $H(z) = h+3$ to $h+4$

Restructuring – Single Rotation


- Rotations to make y the top-most node



- Insertion happens in subtree T_2
- y remains balanced
 - $H(T_1) = h+1$ or $h+2$ or $h+3$
 - ♦ $h+3$ – y is originally unbalanced
 - ♦ $h+2$ – height of y does not increase – z is balanced
 - ♦ therefore $H(T_1) = h+1$
- $H(y) = h+2$ to $h+3$.
- z is unbalanced
 - $H(T_0) = h+1$ or $h+2$ or $h+3$
 - ♦ Since originally z was balanced
 - ♦ $H(T_0) = h+1$
- $H(z) = h+3$ to $h+4$


Restructuring – Single Rotation

- Rotations to make y the top-most node




Restructuring – Single Rotation

- Rotations to make y the top-most node




Restructuring – Single Rotation

- Rotations to make y the top-most node





Restructuring – Single Rotation

- Rotations to make y the top-most node




- Symmetric to the previous scenario



Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.



- Insertion happens in subtree T_1
- $H(T_1) = h$ to $h+1$
- x is balanced
 - $H(T_2) = h$ or $h+1$ or $h+2$
 - ◆ $h+2$ – x is originally unbalanced
 - ◆ $h+1$ – no increase in height of x – z remains balanced
 - ◆ therefore $H(T_2) = h$
- $H(x) = h+1$ to $h+2$


Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Insertion happens in subtree T_1
- y remains balanced

- $H(T_3) = h+1$ or $h+2$ or $h+3$


- ◆ $h+3$ – y is originally unbalanced
- ◆ $h+2$ – no increase in height of $y - z$ remains balanced
- ◆ therefore $H(T_3) = h+1$

- $H(y) = h+2$ to $h+3$



Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Insertion happens in subtree T_1
- y remains balanced
 - $H(T_3) = h+1$ or $h+2$ or $h+3$
 - ◆ $h+3$ – y is originally unbalanced
 - ◆ $h+2$ – no increase in height of y – z remains balanced
 - ◆ therefore $H(T_3) = h+1$
- $H(y) = h+2$ to $h+3$
- z is unbalanced
 - $H(T_0) = h+1$ or $h+2$ or $h+3$
 - ◆ since z was originally balanced
 - ◆ $H(T_0) = h+1$
- $H(z) = h+3$ to $h+4$



Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Insertion happens in subtree T_1
- y remains balanced

- $H(T_3) = h+1$ or $h+2$ or $h+3$

- ◆ $h+3$ – y is originally unbalanced
- ◆ $h+2$ – no increase in height of y – z remains balanced
- ◆ therefore $H(T_3) = h+1$


- $H(y) = h+2$ to $h+3$

- z is unbalanced

- $H(T_0) = h+1$ or $h+2$ or $h+3$


- ◆ since z was originally balanced
- ◆ $H(T_0) = h+1$

- $H(z) = h+3$ to $h+4$




Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- First step
 - rotate x and y




Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- First step
 - rotate x and y




Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Second step
 - rotate x and z




Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Second step
 - rotate x and z




Restructuring – Double Rotation

- Perform rotations around (x,y) and (x,z) to make x the top-most node.


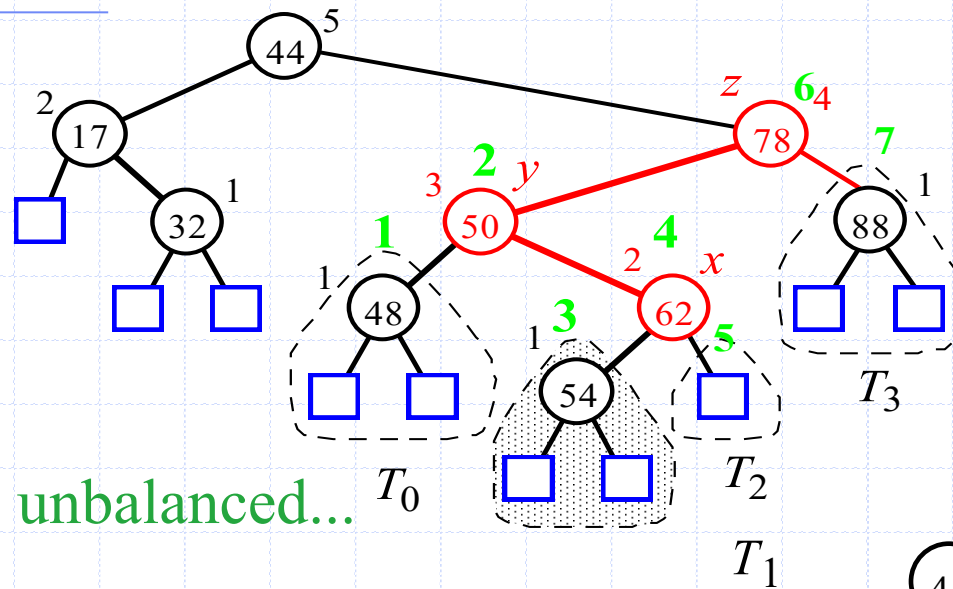


Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- symmetric to the previous configuration




Insertion Example, continued




Removal

- ❑ Removal begins as in a binary search tree, which means the node removed will become an empty external node. Its parent, w , may cause an imbalance.
- ❑ Example:




Rebalancing after a Removal

- Let z be the **first unbalanced** node encountered while travelling up the tree from w . Also, let y be the child of z with the larger height, and let x be the child of y with the larger height
- We perform a **trinode restructuring** to restore balance at z
- As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of T is reached




Deletion



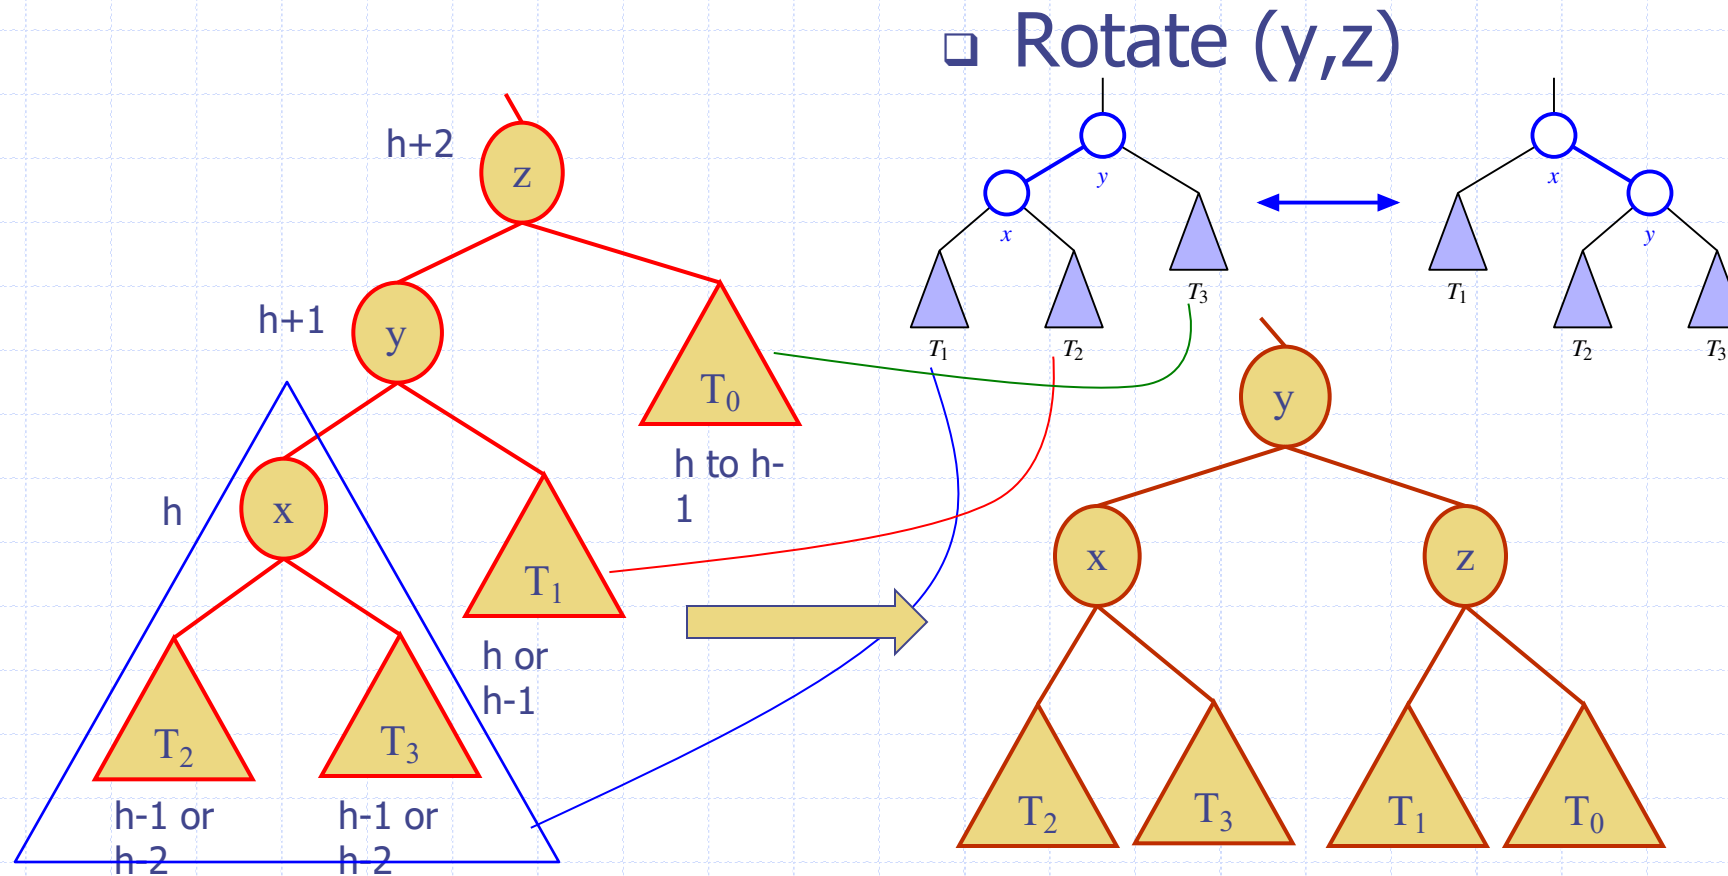
- ❑ Suppose deletion happens in subtree T_0 and its height reduces from h to $h-1$
- ❑ z was originally balanced and now unbalanced
 - $H(y) = h+1$
 - $H(z) = h+2$
- ❑ x has larger height than T_1
 - $H(x) = h$
- ❑ y is balanced
 - $H(T_1) = h$ or $h-1$

Deletion




- Suppose deletion happens in subtree T_0 and its height reduces from h to $h-1$
- x is balanced
 - $H(T_2), H(T_3)$ is $h-1$ or $h-2$
 - However both T_2 and T_3 cannot have height $h-2$

Deletion – Restructuring – Single Rotation




Deletion – Restructuring – Single Rotation

- After rotation the height of the subtree might be less than 1




Deletion – Restructuring – Double Rotation



- Suppose deletion happens in subtree T_0 and its height reduces from h to $h-1$
- z was originally balanced and now unbalanced
 - $H(y) = h+1$
 - $H(z) = h+2$
- x has larger height than T_3
 - $H(x) = h$
- y is balanced
 - $H(T_3) = h-1$


Deletion – Restructuring – Double Rotation




- Suppose deletion happens in subtree T_0 and its height reduces from h to $h-1$
- x remains balanced
 - $H(T_2), H(T_3)$ is $h-1$ or $h-2$
 - However both T_2 and T_3 cannot have height $h-2$

Deletion – Restructuring – Double Rotation

□ Step 1




□ Rotate (x, y)





Deletion – Restructuring – Double Rotation

□ Step 1




□ Rotate (x,y)




Deletion – Restructuring – Double Rotation

□ Step 2




Deletion – Restructuring – Double Rotation

□ Step 2




□ Rotate (x, z)

□ After rotation the height of the subtree is less than 1



Rebalancing after a Removal



AVL Tree Performance

- AVL tree storing n items
 - The data structure uses $O(n)$ space
 - A single restructuring takes $O(1)$ time
 - ◆ using a linked-structure binary tree
 - Searching takes $O(\log n)$ time
 - ◆ height of tree is $O(\log n)$, no restructures needed
 - Insertion takes $O(\log n)$ time
 - ◆ initial find is $O(\log n)$
 - ◆ restructuring up the tree, maintaining heights is $O(\log n)$
 - Removal takes $O(\log n)$ time
 - ◆ initial find is $O(\log n)$
 - ◆ restructuring up the tree, maintaining heights is $O(\log n)$

